

25 evaluation-eligible domains

Each takes one turn as the held-out test set

2 training-only domains

POPS · Real-Life Trial
Always in training — never held out

For each of the 25 evaluation-eligible domains:

One fold (example: "li et al 2015 hotel" held out)

Training set

24 eval-eligible + 2 training-only
= 26 domains

Test set

li et al 2015 hotel
1 held-out domain

Three representations fitted independently on the same split:

TF-IDF

+ Logistic Regression
bag-of-words

AUC · macro-F1

FastText

+ Logistic Regression
static word embeddings

AUC · macro-F1

SBERT

+ Logistic Regression
contextual sentence embeddings

AUC · macro-F1

After all 25 folds:

Results matrix: 25×3

25 folds $\times$ 3 representations
= 75 AUC scores + 75 macro-F1 scores

Repeat for all 25 domains (25 folds total)

 Evaluation-eligible training data

 Held-out test domain

 Fixed training-only (never tested)